


AI Product Life Cycle

Iteration is key!!



Development Tips

Inferencing

- [SEA-LION API](#) available, but limited to 10RPM
- Participants to host their own SEA-LION model on AWS/GCP. Guides are available at <https://docs.sea-lion.ai/guides/inferencing>
- [Ollama](#) for rapid local development

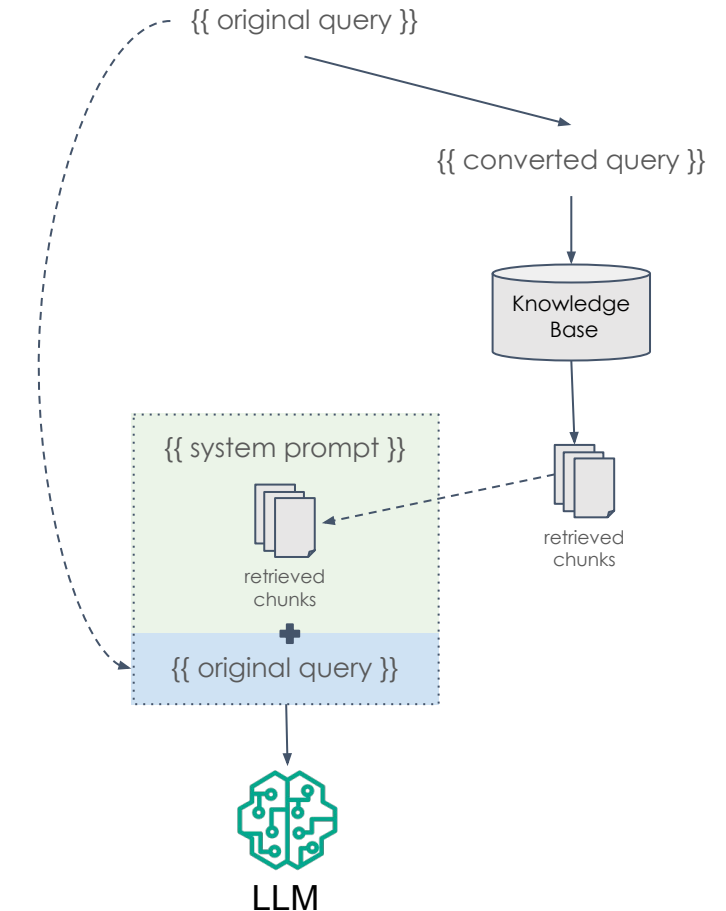
Prompting

- Usual prompting best practices apply
- Can use [Playground Sandbox](#) for prompt testing and model response comparison

Development Tips

RAG

- Usual RAG practices apply
- Embedding model: [BAAI/bge-m3](#)
- Multi-language RAG: what if my user's query is in a different language than my knowledge base (KB) documents?
 - Sending RAG query in a different language would result in poor retrieval of relevant chunks
 - Convert query/extract keywords from query to the language of your KB documents before retrieval
 - Add the retrieved chunks as context to the system prompt, but using the original user query to the LLM. Specify that response is to be generated in the same language as user query
 - This is so that the model responds naturally to the question instead of responding to a bunch of keywords

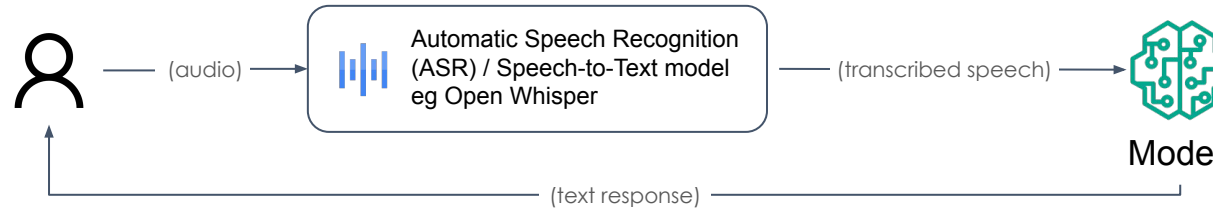


AI SINGAPORE®

Development Tips

Voice

- Pipeline approach



Agents

- Guides available on <https://docs.sea-lion.ai/guides/agents>



Thank you

www.aisingapore.org